

# Oxford: Open-ended learning in symmetric zero-sum games

David Balduzzi w/ Marta Garnelo, Yoram Bachrach, Wojtek  
Czarnecki, Julien Perolat, Max Jaderberg, Thore Graepel



The background of the slide is a dark, textured surface with a complex network of thin, light-colored lines and small dots, resembling a molecular structure or a data network. The lines are interconnected, forming various geometric shapes and patterns. The dots are small and scattered throughout the network, some acting as nodes where multiple lines converge. The overall effect is a sense of depth and complexity, with the network appearing to recede into the distance.

Context



# Peak deep learning

*“If you have a large big  
dataset, and you train a very  
big neural network, then  
success is guaranteed!”*

**Ilya Sutskever (NIPS 2014)**

# Peak deep learning

*"If you have a large big dataset, and you train a very big neural network, then success is guaranteed!"*

**Ilya Sutskever (NIPS 2014)**






# Peak deep learning

*"If you have a large big dataset, and you train a very big neural network, then success is guaranteed!"*

**Ilya Sutskever (NIPS 2014)**

*"If you formulate the right objective, and have enough capacity and compute, then success is guaranteed!"*

**Ilya's law**



# Long ago, not so far away (mid-1800s, Cambridge):



First tutor: **"I'm teaching the most brilliant boy in Britain"**

Second tutor: **"Well, I'm teaching the best test-taker"**

Depending on the version of the story, the first boy was either **Lord Kelvin** or **James Clerk Maxwell**. The second boy indeed scored highest on the Mathematical Tripos, but is otherwise long forgotten.





Long ago, not so far away (mid-1800s, Cambridge):



First tutor: **"I'm teaching the most brilliant boy in Britain"**

Second tutor: **"Well, I'm teaching the best test-taker"**

Depending on the version of the story, the first boy was either **Lord Kelvin** or **James Clerk Maxwell**. The second boy indeed scored highest on the Mathematical Tripos, but is otherwise long forgotten.



**Modern learning algorithms are outstanding test-takers**

Intelligence is about **more than taking tests**  
It's also about **formulating useful problems**



# The problem problem<sup>\*</sup>

**Where do ~~problems~~ objectives come from?**



# Where do problems come from?

## **Answer #1:**

*Someone* packages a dataset into a loss function

e.g. ImageNet, CIFAR, MNIST



# Where do problems come from?

## **Answer #1:**

*Someone* packages a dataset into a loss function

e.g. ImageNet, CIFAR, MNIST

## **Answer #2:**

*Someone* builds a task (that is, an environment sprinkled with rewards)

e.g. Arcade Learning Environment, DM-Lab, Open AI gym




# Where do problems come from?

## Answer #3:

Self-play in symmetric zero-sum games

The agent *is* the task -- create an outer loop that applies deep RL to itself





# (Naive) self-play *is* an open-ended learning algorithm

It's pretty amazing

---

**Algorithm 2** Self-play

---

**input:** agent  $\mathbf{v}_1$   
**for**  $t = 1, \dots, T$  **do**  
     $\mathbf{v}_{t+1} \leftarrow \text{oracle}(\mathbf{v}_t, \phi_{\mathbf{v}_t}(\bullet))$   
**end for**  
**output:**  $\mathbf{v}_{T+1}$

---






# (Naive) self-play *is* an open-ended learning algorithm

but ...

there are really simple examples  
where it completely breaks down

It's **not** a general purpose learning  
algorithm, *not even* for zero-sum games





# Open questions

## 1. When does self-play work, and why?

- If self-play is the answer then what is the question?
- “works for some games but not others” isn't good enough

## 2. What is the goal in (general) zero-sum games?

- Yes, to win, but against whom? The opponent matters
  - “Beating a pro” isn't a formal specification
- Is there a general algorithm?



# What does success look like?

**#1:** Supervised learning (e.g. ImageNet)

**#2:** Reinforcement learning on fixed environment (e.g. DM-Lab)

**#3:** Self-play (e.g. on Chess, Go, Shogi, ... ; but **not** on everything )



# What underlies machine learning's big wins?

## 1. A transitive performance measure

- e.g. a loss function or discounted rewards

## 2. An incremental improvement operator

- e.g. gradient descent, RL, evolutionary algorithms



# What underlies machine learning's big wins?

## 1. A transitive performance measure

- e.g. a loss function or discounted rewards

## 2. An incremental improvement operator

- e.g. gradient descent, RL, evolutionary algorithms

What drives **self-play**'s success?

1. **Improvement (local):**  $(A_{t+1} > A_t)$  and  $(A_t > A_{t-1})$  and ...
2. **Transitivity (global):**  $A_{\text{final}}$  beats all previous agents

The background of the slide is a dark, textured surface with a complex network of thin, light-colored lines and small dots, resembling a molecular structure or a data network. The lines are interconnected, forming various geometric shapes and patterns. The dots are small and light-colored, scattered throughout the network. The overall effect is a sense of depth and complexity.

# Basics



# On the varieties of zero-sum games



**transitive:** “relative skill determines who wins”

**cyclic:** “every strategy has a counter-strategy”



# Convex → Nonconvex → Transitive → Nontransitive

**Convex objectives** (ML in early 2000s):  
Converge to **unique global optimum**, for  
any initial condition and for any  
reasonable algorithm.

**Nonconvex objectives** (deep learning,  
~2010...): Converge to a **local optimum**  
that depends on choice of initial  
condition and choice of algorithm.

**Transitive games** (e.g. self-play in Go,  
2016...): Optimize and generate a  
sequence of objectives of the **same type**  
and **increasing difficulty**. Best agent at  
hardest objective is best overall.

**Nontransitive games** (e.g. AlphaStar  
league, 2019...): Optimum depends on  
who you compete with. Rather than  
finding a single best agent, should **invest**  
**in training a diverse “ecology” of agents.**



**Functional-form game:** a two-player symmetric zero-sum game, with differentiable parametrization:

$$\phi(\mathbf{v}, \mathbf{w})$$

**Note:** function approximator is “folded into” the game's definition.





**Theorem:** Any symmetric zero-sum game decomposes into

[ **transitive** ] + [ **cyclic** ] components


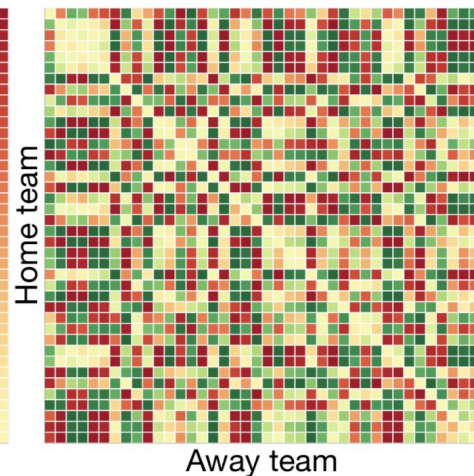
**transitive:** skill determines  
outcome

**cyclic:** every strategy has a  
counter-strategy

Transitive game



Cyclic game








# Transitive and cyclic games

	Good	Better	Best	Avg
Good	0	-1	-2	-1.5
Better	1	0	-1	0.0
Best	2	1	0	1.5


	Rock	Paper	Scissors	Avg
Rock	0	-1	1	0.0
Paper	1	0	-1	0.0
Scissors	-1	1	0	0.0



Transitive game



Cyclic game





# Transitive games: where the opponent doesn't matter

$$\phi(\mathbf{v}, \mathbf{w}) = f(\mathbf{v}) - f(\mathbf{w})$$

- $f$  assigns **rating** to players
- outcome is difference in ratings

→ algorithm: **optimize against a fixed opponent**

$$\operatorname{argmax}_{\mathbf{v}} \phi(\mathbf{v}, \mathbf{w})$$

for any, fixed,  $\mathbf{w}$



# Monotonic games: what self-play is “meant for”

$$\phi(\mathbf{v}, \mathbf{w}) = \sigma( f(\mathbf{v}) - f(\mathbf{w}) )$$

- $\sigma$  is a monotonic function (e.g. **Elo's rating model**)
- no learning signal if you train against a weak opponent (gradients vanish)
  - optimizing against a fixed opponent doesn't work

→ algorithm: **self-play**

---

**Algorithm 2** Self-play

---

**input:** agent  $\mathbf{v}_1$   
**for**  $t = 1, \dots, T$  **do**  
     $\mathbf{v}_{t+1} \leftarrow \text{oracle}(\mathbf{v}_t, \phi_{\mathbf{v}_t}(\bullet))$   
**end for**  
**output:**  $\mathbf{v}_{T+1}$

---



# Cyclic games: where naive self-play breaks down

A game is cyclic if  $\int_W \phi(\mathbf{v}, \mathbf{w}) \cdot d\mathbf{w} = 0$


i.e **wins, against some opponents, are balanced by losses, against others.**

Implications:

- There is no best agent
- Measuring performance of individuals is nonsensical
  - which is better, paper or rock?



NEED: a **transitive objective**,  
in **non-transitive games**!



# How to convert **Agents** $\rightarrow$ **Objectives**

**Definition:** **Gamescape** is the convex hull of all objectives in a game

A game is a function that evaluates pairs of agents:

$$\Phi : W \times W \rightarrow R$$

$\Phi(\mathbf{v}, \mathbf{w})$  (e.g. probability that  $\mathbf{v}$  beats  $\mathbf{w}$ )

**Fixing an opponent** converts an **Agent**  $\rightarrow$  **Objective function**

$$\mathbf{w} \rightarrow \Phi_{\mathbf{w}}(-) := \Phi(-, \mathbf{w})$$

currying!



# Gamescapes

**Functional Gamescape:** convex hull of all objective functions  $\Phi_w(-)$

- lives in function space  $\mathcal{F}(W, R)$
- intractable

**Empirical Gamescape:** convex mixture of all rows of evaluation matrix

- proxy for functional gamescape
- tractable



**Fitness landscapes** are a special case that arise when the game is transitive or monotone.



# Landscape

**Transitive:** there's **one objective**  
("improve skill") ...



... so gamescape degenerates to  
**one dimensional fitness  
landscape**



# Gamescape

**Nontransitive:** different opponents  
are **different objectives** that pull in  
different directions


Gamescape is **multi-dimensional**






# Gamescapes


Almost Transitive








Mixed



Almost Cyclic



Random





mElo / Schur decomposition (roughly, in rank-2 case)

$$A_{n \times n} = W_{n \times 2} \cdot \begin{pmatrix} 0 & 1 \\ -1 & 0 \end{pmatrix} \cdot W_{2 \times n}^T$$

{ agents }  $\rightarrow$  { rows of **W** }

Entries of each row are the agent's **multi-dimensional Elo** scores  
(let's forget about sigmoids)



# Loooooooooong cycles

Is transitive geometry always 1-dim? **Yes**


Is cyclic geometry always 2-dim? **No**

**Rock** → **Paper** → **Scissors** → **Fire** → **Water** → **Air** → **Ether** → **Milk** → ... → **Rock**

**length(cycle)** = {  $2n$  or  $2n+1$  } → **rank(A)** = {  $2n-2$  or  $2n$  }, respectively.

What is the dimension of SC2's gamescape? No idea.

Ceci n'est pas une pipe





So what's **the transitive objective** in nontransitive games?

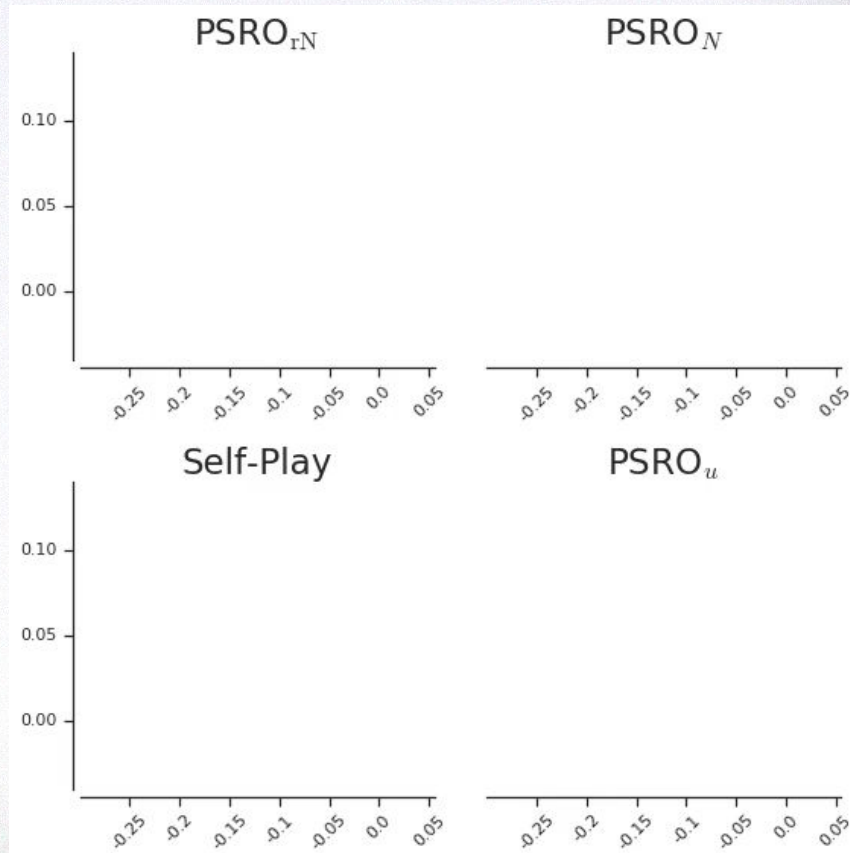
**Individual** performance is **meaningless** in nontransitive games.

“Which is better, rock or paper?”

Idea: **population-level objectives!**

**Grow the gamescape**, to get  
“diverse, effective agents”

- find strategic dimensions of game
- and best ways of executing them



# Population-level performance

## Definition:

Given  $(m \times n)$  evaluation matrix  $A$  for populations  $P$  and  $Q$ .

Let  $(p, q)$  be any Nash Equilibrium on zero-sum meta-game on  $A$ .

## Relative performance of populations:

$$v(P, Q) := p^T \cdot A \cdot q$$

(doesn't depend on choice of Nash).

## Intuition:

- I pick my mixture of champions
- You pick yours (simultaneously)
- We play them out and
- See who wins, how much, in expectation



# Population-level performance


**Lemma (**transitivity** for populations):**

If agents in  $P$  are convex combinations of agents in  $Q$  then

$$v(P, Q) \leq 0$$

**Implication:**

Algorithms that “grow polytope” are guaranteed to improve the population-level performance.





# Two algorithms: one old and one new

Two ways to **track growth** ...

1. Population performance
2. Effective diversity

... and two algorithms

1. Response to Nash (double oracle)
  - train against Nash
2. Response to **rectified** Nash
  - train against agents in Nash that you beat
  - game-theoretic niches

---

**Algorithm 3** Response to Nash (PSRO<sub>N</sub>)

---

**input:** population  $\mathfrak{P}_1$  of agents  
**for**  $t = 1, \dots, T$  **do**  
     $\mathbf{p}_t \leftarrow \text{Nash on } \mathbf{A}_{\mathfrak{P}_t}$   
     $\mathbf{v}_{t+1} \leftarrow \text{oracle}(\mathbf{v}_t, \sum_{\mathbf{w}_i \in \mathfrak{P}_t} \mathbf{p}_t[i] \cdot \phi_{\mathbf{w}_i}(\bullet))$   
     $\mathfrak{P}_{t+1} \leftarrow \mathfrak{P}_t \cup \{\mathbf{v}_{t+1}\}$   
**end for**  
**output:**  $\mathfrak{P}_{T+1}$

---

---

**Algorithm 4** Response to rectified Nash (PSRO<sub>rN</sub>)

---


**input:** population  $\mathfrak{P}_1$   
**for**  $t = 1, \dots, T$  **do**  
     $\mathbf{p}_t \leftarrow \text{Nash on } \mathbf{A}_{\mathfrak{P}_t}$   
    **for** agent  $\mathbf{v}_t$  with positive mass in  $\mathbf{p}_t$  **do**  
         $\mathbf{v}_{t+1} \leftarrow \text{oracle}(\mathbf{v}_t, \sum_{\mathbf{w}_i \in \mathfrak{P}_t} \mathbf{p}_t[i] \cdot \lfloor \phi_{\mathbf{w}_i}(\bullet) \rfloor_+)$   
    **end for**  
     $\mathfrak{P}_{t+1} \leftarrow \mathfrak{P}_t \cup \{\mathbf{v}_{t+1} : \text{updated above}\}$   
**end for**  
**output:**  $\mathfrak{P}_{T+1}$

---




# Response to Rectified Nash


**A**



**B**

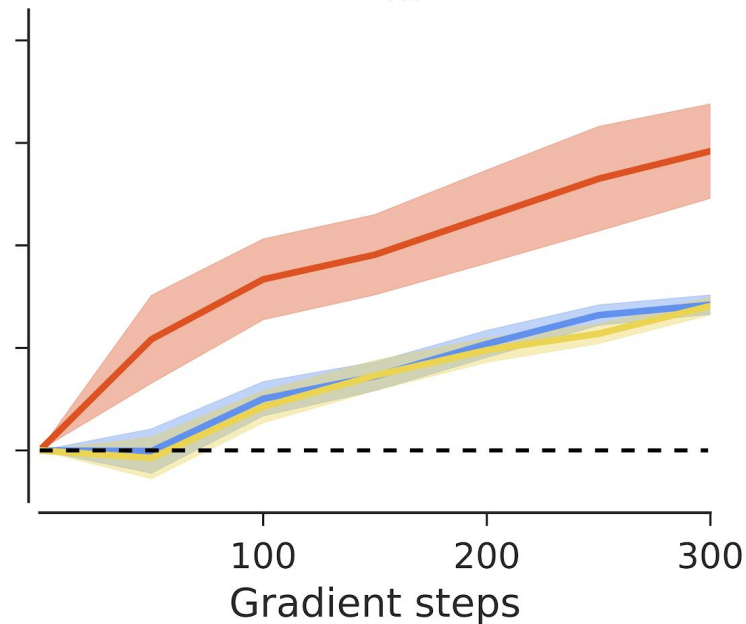
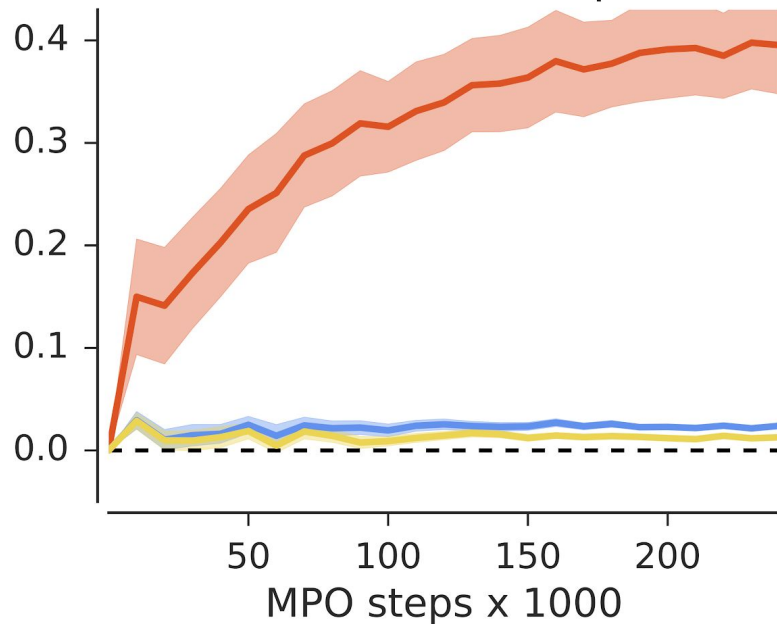


**C**



# Performance on Blotto and differentiable Lotto

Relative Population Performance of  $\text{PSRO}_{rN}$




—  $\text{PSRO}_{rN}$  vs. Self-play

—  $\text{PSRO}_{rN}$  vs.  $\text{PSRO}_U$

—  $\text{PSRO}_{rN}$  vs.  $\text{PSRO}_N$



# Growth of gamescapes in differentiable Lotto




The background of the slide is a dark gray to black gradient. It features a complex, abstract network of thin, light gray lines connecting various points. Some points are represented by small, solid gray circles, while others are just intersections of lines. The network is dense and irregular, with some areas having more connections than others, creating a sense of depth and complexity. The overall effect is reminiscent of a molecular structure or a data network visualization.

# Discussion




Back to the problem problem





## Back to the problem problem

1. **How can algorithms formulate problems?**
2. **Which problems are useful?**





# Back to the problem problem


1. How can algorithms formulate problems?
2. Which problems are useful?

**Formulating problems** is an **inter-agent** phenomenon

- in a zero-sum game, “my behavior is your problem”

**Usefulness** is **semantically grounded** in **winning** in ZSGs


- so **Nash formulates useful problems** via, say, response-to-Nash, response-to-rectified-Nash





# Comment on human problem posing

- The model describes posing-and-solving problems as a population-level activity
- This is how it works for humans
  - At least, population aspects are under-rated;
  - individuals get credit for work of populations
- Even humans like Einstein didn't pose new problems
  - "Everyone" knew Maxwell's equations aren't invariant under Galilean transforms
  - Einstein "just" took the problem seriously





# Summary

Success of modern machine boils down to

- applying many **incremental improvements**
- to **transitive objectives**, so improvements **add up**

This talk:

- tools for **formulating objectives in a nontransitive setting**
- and **algorithms for optimizing them**

details: <https://arxiv.org/abs/1901.08106>